

# Used Vehicles Analysis and Price Prediction

- About the data set:
  - This data set is used vehicles data publically available on <https://www.kaggle.com/datasets/austinreese/craigslist-carstrucks-data>
- Description of data variables:
  - 'year': model year of a vehicle | int32
  - 'manufacturer': maker of a vehicle | categorical
  - 'cylinder': cylinder size of a vehicle | categorical
  - 'type': body type of a vehicle | categorical
  - 'paint\_color': color of a vehicle | categorical
  - 'region': region where the vehicle was listed | categorical
  - 'state': state where the vehicle was listed | categorical
  - 'price': listing price on Craigslist.org | float
- In descriptive analysis, I explore which feature group(s) of vehicles have greater impact on listing prices, then followed with a correlation analysis of numerical variables (including converted categorical variable to numeric variable).
- In inferential analysis, I use the linear regression machine learning method to build five models, selecting the most accurate one to predict the price of used vehicles as desired.

## Part 1: Import necessary libraries, load data, and check data types and structures

```
In [20]: # Import necessary libraries:
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

```
import pandas as pd
import statsmodels.api as sm
import statsmodels.formula.api as smf
from statsmodels.formula.api import ols
```

```
In [21]: # Load dataset:
df = pd.read_csv('vehicles.csv', header=0)

# remove white space from cell and header
df = df.apply(lambda x: x.astype(str).str.strip() if x.dtype == "object" else x)
df.columns=df.columns.str.strip()
```

```
In [22]: # check column names, data types, total observations:
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 426880 entries, 0 to 426879
Data columns (total 26 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   id                    426880 non-null  int64
 1   url                   426880 non-null  object
 2   region                426880 non-null  object
 3   region_url            426880 non-null  object
 4   price                 426880 non-null  int64
 5   year                  425675 non-null  float64
 6   manufacturer          426880 non-null  object
 7   model                 426880 non-null  object
 8   condition             426880 non-null  object
 9   cylinders             426880 non-null  object
10   fuel                  426880 non-null  object
11   odometer              422480 non-null  float64
12   title_status          426880 non-null  object
13   transmission          426880 non-null  object
14   VIN                   426880 non-null  object
15   drive                 426880 non-null  object
16   size                  426880 non-null  object
17   type                  426880 non-null  object
18   paint_color           426880 non-null  object
19   image_url             426880 non-null  object
20   description           426880 non-null  object
21   county                0 non-null       float64
22   state                 426880 non-null  object
23   lat                   420331 non-null  float64
24   long                  420331 non-null  float64
25   posting_date          426880 non-null  object
dtypes: float64(5), int64(2), object(19)
memory usage: 84.7+ MB

```

```

In [127... # A List of vehicles' model year for sale:
df.year.unique()

```

```
Out[127... array([ nan, 2014., 2010., 2020., 2017., 2013., 2012., 2016., 2019.,
        2011., 1992., 2018., 2004., 2015., 2001., 2006., 1968., 2003.,
        2008., 2007., 2005., 1966., 2009., 1998., 2002., 1999., 2021.,
        1997., 1976., 1969., 1995., 1978., 1954., 1979., 1970., 1974.,
        1996., 1987., 2000., 1955., 1960., 1991., 1972., 1988., 1994.,
        1929., 1984., 1986., 1989., 1973., 1946., 1933., 1958., 1937.,
        1985., 1957., 1953., 1942., 1963., 1977., 1993., 1903., 1990.,
        1965., 1982., 1948., 1983., 1936., 1932., 1951., 1931., 1980.,
        1967., 1971., 1947., 1981., 1926., 1962., 1975., 1964., 1934.,
        1952., 1940., 1959., 1950., 1930., 1956., 1922., 1928., 2022.,
        1901., 1941., 1924., 1927., 1939., 1923., 1949., 1961., 1935.,
        1918., 1900., 1938., 1913., 1916., 1943., 1925., 1921., 1915.,
        1945., 1902., 1905., 1920., 1944., 1910., 1909.]])
```

```
In [129... # check null values:
df.isna().sum()
```

```
Out[129... id          0
          url          0
          region       0
          region_url   0
          price        0
          year         1205
          manufacturer  0
          model         0
          condition     0
          cylinders     0
          fuel          0
          odometer     4400
          title_status  0
          transmission  0
          VIN          0
          drive        0
          size         0
          type         0
          paint_color   0
          image_url     0
          description   0
          county       426880
          state        0
          lat          6549
          long         6549
          posting_date  0
          dtype: int64
```

## Part 2: Preprocess the data

```
In [131... # Drop columns that won't be useful for this analysis:
df.drop(columns=['id','url','region_url','model','condition','VIN','image_url','county','transmission','title_status'])
```

```
In [132... # Remove rows that impact accuracy of the analysis:
df= df[(df.price >= 2000) & (df.price <=500000) & (df.year >=1970)& (df.year < 2022) &(df.odometer<300000) & (df.odor
```

```
In [133... # Remove null values in the cells:
df=df.dropna(subset=['lat','long','drive','lat','long'])
```

```
In [152... # Remove 'nan' strings in those cells:
df=df[(df.year!='nan')&(df.price !='nan')&(df.manufacturer !='nan')&(df.cylinders !='nan')&(df.fuel !='nan')&(df.odor

# Recheck total observations:
df.shape
```

Out[152... (137469, 13)

```
In [154... # Change data type for year column:
df.year=df.year.astype(int)
```

## Part 3: Descriptive statistics analysis

```
In [157... # re-check data structure and data types:
df.info()
```

<class 'pandas.core.frame.DataFrame'>  
Index: 137469 entries, 31 to 426878  
Data columns (total 13 columns):

#	Column	Non-Null Count	Dtype
0	region	137469 non-null	object
1	price	137469 non-null	int64
2	year	137469 non-null	int32
3	manufacturer	137469 non-null	object
4	cylinders	137469 non-null	object
5	fuel	137469 non-null	object
6	odometer	137469 non-null	float64
7	drive	137469 non-null	object
8	type	137469 non-null	object
9	paint_color	137469 non-null	object
10	state	137469 non-null	object
11	lat	137469 non-null	float64
12	long	137469 non-null	float64

dtypes: float64(3), int32(1), int64(1), object(8)  
memory usage: 14.2+ MB

```
In [159... # Re-check missing values:
df.isnull().sum()
```


```
Out[159... region      0
price      0
year       0
manufacturer 0
cylinders  0
fuel       0
odometer   0
drive      0
type       0
paint_color 0
state      0
lat        0
long       0
dtype: int64
```

```
In [161... # re-check statistical summary:
df.describe()
```

```
Out[161...
```

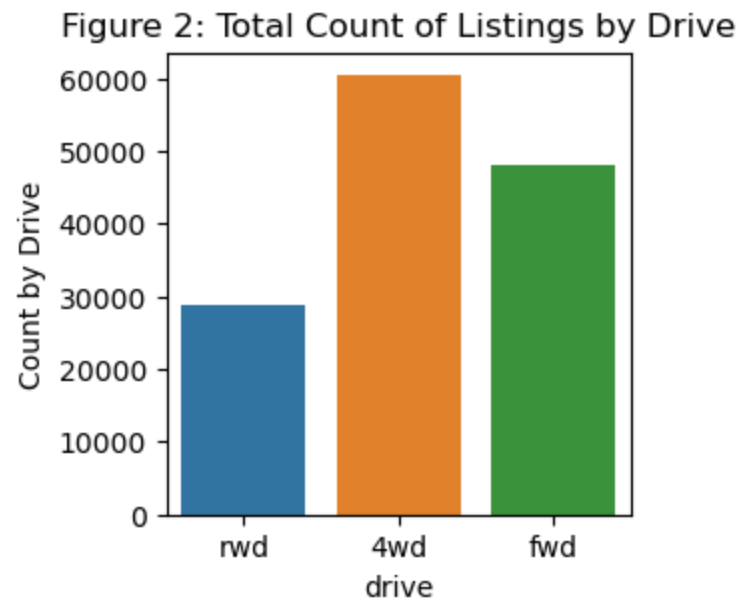
	price	year	odometer	lat	long
<b>count</b>	137469.000000	137469.000000	137469.000000	137469.000000	137469.000000
<b>mean</b>	17773.750198	2011.044701	101588.451171	38.634585	-94.335534
<b>std</b>	13732.483014	7.051118	59477.895554	5.984344	18.450910
<b>min</b>	2000.000000	1970.000000	501.000000	-81.838232	-159.719900
<b>25%</b>	7400.000000	2008.000000	53000.000000	34.969253	-110.801411
<b>50%</b>	13800.000000	2013.000000	99095.000000	39.444084	-88.105874
<b>75%</b>	25590.000000	2016.000000	142000.000000	42.402459	-80.406800
<b>max</b>	289995.000000	2021.000000	299999.000000	82.252826	139.691700

```
In [163... # count by color group:
plt.figure(figsize=(15,3))
sns.countplot(data=df, x='paint_color', hue='paint_color')
plt.title('Figure 1: Total Count of Listings by Color')
plt.ylabel('Count by Color')
plt.show()
```



In [274...

```
# count by drive group:
plt.figure(figsize=(3,3))
sns.countplot(data=df, x='drive', hue='drive')
plt.title('Figure 2: Total Count of Listings by Drive')
plt.ylabel('Count by Drive')
plt.show()
```




In [167...


```
# count by manufacturer group:
plt.figure(figsize=(20,3))
```



```
sns.countplot(data=df, x='manufacturer', hue='manufacturer')
plt.title('Figure 2: Total Count of Listings by Manufacturer')
plt.ylabel('Count by Manufacturer')
plt.xticks(rotation=45)
plt.show()
```



```
In [173... # count by year group:
plt.figure(figsize=(20,3))
sns.countplot(data=df, x='year')
plt.title('Figure 3: Total Count of Listings by Year')
plt.ylabel('Count by Year')
plt.xticks(rotation=45)
plt.show()
```



```
In [175... # save the file to a csv file after cleaning for further visulization analysis using Tableau:
df.to_csv('vehicles_clean.csv', index= False)
```

Part 3 conclusions:

- White color of vehicle has the most number of listings, 4wd - Four wheel drive is the most popular drive type, followed by fwd - front wheel drive)

## Part 4: Correlation analysis

In order to find out whether there are correlations between price and other non-numerical variables such as drive, color, cylinders, we will need to replace categorical name to numerical numbers for the pupose of correlation analysis.

- To replace categorical names to numerical numbers, for example:

`drive['rwd', '4wd','fwd'] -- replaced with -- drive_n [1, 3, 2]`

```
In [180... dfn=pd.read_csv('vehicles_clean.csv', header=0)
```

```
In [182... dfn
```

Out[182...

	region	price	year	manufacturer	cylinders	fuel	odometer	drive	type	paint_color	state	lat	lon
<b>0</b>	auburn	15000	2013	ford	6 cylinders	gas	128000.0	rwd	truck	black	al	32.592000	-85.518900
<b>1</b>	auburn	27990	2012	gmc	8 cylinders	gas	68696.0	4wd	pickup	black	al	32.590000	-85.480000
<b>2</b>	auburn	34590	2016	chevrolet	6 cylinders	gas	29499.0	4wd	pickup	silver	al	32.590000	-85.480000
<b>3</b>	auburn	35000	2019	toyota	6 cylinders	gas	43000.0	4wd	truck	grey	al	32.601300	-85.443970
<b>4</b>	auburn	29990	2016	chevrolet	6 cylinders	gas	17302.0	4wd	pickup	red	al	32.590000	-85.480000
...	...	...	...	...	...	...	...	...	...	...	...	...	...
<b>137464</b>	wyoming	48590	2020	cadillac	6 cylinders	gas	7701.0	fwd	other	black	wy	33.779214	-84.411810
<b>137465</b>	wyoming	39990	2017	infiniti	8 cylinders	gas	41664.0	4wd	other	black	wy	33.779214	-84.411810
<b>137466</b>	wyoming	32990	2016	infiniti	8 cylinders	gas	55612.0	rwd	other	black	wy	33.779214	-84.411810
<b>137467</b>	wyoming	33590	2018	lexus	6 cylinders	gas	30814.0	rwd	sedan	white	wy	33.779214	-84.411810
<b>137468</b>	wyoming	28990	2018	lexus	6 cylinders	gas	30112.0	fwd	sedan	silver	wy	33.786500	-84.445400

137469 rows × 13 columns



In [184...

```
dfn.dtypes
```

```
Out[184... region      object
price       int64
year        int64
manufacturer object
cylinders   object
fuel        object
odometer    float64
drive       object
type        object
paint_color object
state       object
lat         float64
long        float64
dtype: object
```

```
In [186... dfn.describe()
```

```
Out[186...
```

	price	year	odometer	lat	long
<b>count</b>	137469.000000	137469.000000	137469.000000	137469.000000	137469.000000
<b>mean</b>	17773.750198	2011.044701	101588.451171	38.634585	-94.335534
<b>std</b>	13732.483014	7.051118	59477.895554	5.984344	18.450910
<b>min</b>	2000.000000	1970.000000	501.000000	-81.838232	-159.719900
<b>25%</b>	7400.000000	2008.000000	53000.000000	34.969253	-110.801411
<b>50%</b>	13800.000000	2013.000000	99095.000000	39.444084	-88.105874
<b>75%</b>	25590.000000	2016.000000	142000.000000	42.402459	-80.406800
<b>max</b>	289995.000000	2021.000000	299999.000000	82.252826	139.691700

```
In [244... # Remove string from column cylinders:
dfn['cylinders'] = dfn['cylinders'].astype(str)
dfn['cylinders'] = dfn['cylinders'].str.replace(' cylinders', '', regex=False)
dfn['cylinders'] = dfn['cylinders'].str.replace('other', '13', regex=False)
dfn['cylinders'] = dfn['cylinders'].astype(int)
dfn.cylinders.unique()
```

```
Out[244... array([ 6,  8,  4,  5, 10,  3, 13, 12])
```

```
In [256... # add a column age, calculate the age of the vehicle by subtracting model year from current year when the data was c  
dfn['age'] = 2022 - dfn['year']  
dfn.age.unique()
```

```
Out[256... array([ 9, 10,  6,  3, 11,  5,  4, 18, 21,  8, 19, 14, 15,  2, 16, 17, 12,  
        13, 20, 27, 48, 26, 22, 23,  7, 25, 24, 50,  1, 28, 43, 38, 46, 31,  
        29, 32, 44, 40, 36, 34, 37, 33, 41, 30, 49, 47, 51, 35, 42, 45, 52,  
        39], dtype=int64)
```

```
In [258... # Ensure 'drive' is a string column:  
dfn['drive'] = dfn['drive'].astype(str)  
  
# Replace the values using a dictionary:  
dfn['drive'] = dfn['drive'].replace({'rwd': '1', '4wd': '3', 'fwd': '2'})  
  
# Convert 'drive' back to integer:  
dfn['drive'] = dfn['drive'].astype(int)  
  
# Check the unique values in the 'drive' column:  
print(dfn['drive'].unique())
```

```
[1 3 2]
```

```
In [260... dfn=dfn.dropna()
```

```
In [262... dfn.to_csv('clean.csv')
```

## EDA


```
In [264... dfn.head()
```

Out[264...

	region	price	year	manufacturer	cylinders	fuel	odometer	drive	type	paint_color	state	lat	long	age
0	auburn	15000	2013	ford	6	gas	128000.0	1	truck	black	al	32.5920	-85.518900	9
1	auburn	27990	2012	gmc	8	gas	68696.0	3	pickup	black	al	32.5900	-85.480000	10
2	auburn	34590	2016	chevrolet	6	gas	29499.0	3	pickup	silver	al	32.5900	-85.480000	6
3	auburn	35000	2019	toyota	6	gas	43000.0	3	truck	grey	al	32.6013	-85.443974	3
4	auburn	29990	2016	chevrolet	6	gas	17302.0	3	pickup	red	al	32.5900	-85.480000	6


In [270...

```
plt.figure(figsize=(3, 3))
sns.histplot(dfn['age'], kde=False, bins=20)
# Show the plot
plt.show()
```




In [272...

```
plt.figure(figsize=(3, 3))
sns.histplot(dfn['odometer'], kde=False, bins=20)
plt.show()
```




```
In [240... sns.scatterplot(data=dfn, x='price',y='odometer')  
plt.show()  
plt.figure(figsize=(3,3))
```



```
Out[240... <Figure size 300x300 with 0 Axes>
```


```
In [242... sns.scatterplot(data=dfn, x='price',y='age')  
plt.show()  
plt.figure(figsize=(3,3))
```



```
Out[242... <Figure size 300x300 with 0 Axes>
```


```
In [209... sns.scatterplot(data=dfn, x='price',y='drive')  
plt.show()  
plt.figure(figsize=(3,3))
```





Out[209... <Figure size 300x300 with 0 Axes>

```
In [211... sns.scatterplot(data=dfn, x='price',y='cylinders')  
plt.show()  
plt.figure(figsize=(3,3))
```




Out[211... <Figure size 300x300 with 0 Axes>

It seems like there is some negative relationship between price and odometer and age. let's use heatmap to show the strength of the relationship.

```
In [214... # Create a heatmap to visualize the correlations:
corr = dfn[['price', 'age', 'cylinders', 'odometer', 'drive']]
plt.figure(figsize=(12,3))
sns.heatmap(corr.corr(), annot=True, cmap='Reds')
#plt.title('Figure 4: Correlation Heatmap')
plt.show()
```

<Figure size 300x300 with 0 Axes>



#### Part 4 Conclusion:

- Correlation analysis suggested:
  - a negative linear relationship between price(the dependent variable) and the explanatory variables: odometer, age.
  - a positive linear relationship between price and cylinders and drive.

## Part 5: Inferential analysis

```
In [218... model1 = smf.ols("price ~ odometer", data=dfn).fit()

print(model1.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          price    R-squared:                0.312
Model:                  OLS      Adj. R-squared:            0.312
Method:                 Least Squares    F-statistic:          6.237e+04
Date:                   Sun, 23 Feb 2025    Prob (F-statistic):      0.00
Time:                   10:14:53    Log-Likelihood:         -1.4791e+06
No. Observations:       137469    AIC:                   2.958e+06
Df Residuals:           137467    BIC:                   2.958e+06
Df Model:                1
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	3.088e+04	60.800	507.849	0.000	3.08e+04	3.1e+04
odometer	-0.1290	0.001	-249.738	0.000	-0.130	-0.128

```

=====
Omnibus:                 71798.292    Durbin-Watson:           1.542
Prob(Omnibus):           0.000    Jarque-Bera (JB):        1456878.838
Skew:                    2.067    Prob(JB):                 0.00
Kurtosis:                18.403    Cond. No.                 2.33e+05
=====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 2.33e+05. This might indicate that there are
strong multicollinearity or other numerical problems.
```

```
In [220... model2 = smf.ols("price ~ age", data=dfn).fit()

print(model2.summary())
```

# OLS Regression Results

```
=====
Dep. Variable:          price    R-squared:                0.230
Model:                  OLS      Adj. R-squared:            0.230
Method:                 Least Squares    F-statistic:          4.103e+04
Date:                  Sun, 23 Feb 2025    Prob (F-statistic):    0.00
Time:                  10:14:54    Log-Likelihood:       -1.4868e+06
No. Observations:      137469    AIC:                  2.974e+06
Df Residuals:          137467    BIC:                  2.974e+06
Df Model:              1
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.8e+04	60.057	466.269	0.000	2.79e+04	2.81e+04
age	-933.7176	4.610	-202.552	0.000	-942.753	-924.683

```
=====
Omnibus:              75568.150    Durbin-Watson:          1.525
Prob(Omnibus):        0.000    Jarque-Bera (JB):       1425392.091
Skew:                 2.246    Prob(JB):               0.00
Kurtosis:             18.122    Cond. No.:               24.2
=====
```

## Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [222... model3 = smf.ols('price ~ age + odometer', data=dfn).fit()

print(model3.summary())
```

# OLS Regression Results

```

=====
Dep. Variable:          price    R-squared:                0.374
Model:                  OLS      Adj. R-squared:            0.374
Method:                 Least Squares    F-statistic:            4.111e+04
Date:                  Sun, 23 Feb 2025    Prob (F-statistic):      0.00
Time:                  10:14:55    Log-Likelihood:         -1.4726e+06
No. Observations:      137469    AIC:                    2.945e+06
Df Residuals:          137466    BIC:                    2.945e+06
Df Model:              2
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	3.383e+04	63.242	534.862	0.000	3.37e+04	3.39e+04
age	-547.6108	4.687	-116.842	0.000	-556.797	-538.425
odometer	-0.0990	0.001	-178.103	0.000	-0.100	-0.098

```

=====
Omnibus:                79422.138    Durbin-Watson:           1.540
Prob(Omnibus):          0.000    Jarque-Bera (JB):        2119328.570
Skew:                   2.288    Prob(JB):                 0.00
Kurtosis:               21.683    Cond. No.:                2.54e+05
=====

```

## Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.54e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```

In [224... model4 = smf.ols('price ~ age + odometer + cylinders', data=dfn).fit()

print(model4.summary())

```

# OLS Regression Results

```

=====
Dep. Variable:          price    R-squared:                0.553
Model:                  OLS      Adj. R-squared:            0.553
Method:                 Least Squares    F-statistic:          5.673e+04
Date:                   Sun, 23 Feb 2025    Prob (F-statistic):    0.00
Time:                   10:14:56    Log-Likelihood:       -1.4494e+06
No. Observations:      137469    AIC:                  2.899e+06
Df Residuals:          137465    BIC:                  2.899e+06
Df Model:               3
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.415e+04	99.428	142.363	0.000	1.4e+04	1.43e+04
age	-692.7431	4.008	-172.819	0.000	-700.600	-684.887
odometer	-0.0965	0.000	-205.526	0.000	-0.097	-0.096
cylinders	3534.9838	15.067	234.611	0.000	3505.452	3564.516

```

=====
Omnibus:                94577.747    Durbin-Watson:          1.673
Prob(Omnibus):           0.000    Jarque-Bera (JB):       5407315.890
Skew:                    2.691    Prob(JB):               0.00
Kurtosis:                33.250    Cond. No.:              4.77e+05
=====

```

## Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 4.77e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```

In [226... model5 = smf.ols('price ~ age + odometer + drive + cylinders', data=dfn).fit()

print(model5.summary())

```

### OLS Regression Results

```

=====
Dep. Variable:          price    R-squared:                0.568
Model:                  OLS      Adj. R-squared:           0.568
Method:                 Least Squares    F-statistic:           4.511e+04
Date:                   Sun, 23 Feb 2025    Prob (F-statistic):      0.00
Time:                   10:14:56    Log-Likelihood:         -1.4472e+06
No. Observations:       137469    AIC:                    2.894e+06
Df Residuals:           137464    BIC:                    2.894e+06
Df Model:                4
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	9361.4145	120.706	77.556	0.000	9124.833	9597.996
age	-632.4618	4.042	-156.462	0.000	-640.385	-624.539
odometer	-0.1016	0.000	-217.117	0.000	-0.103	-0.101
drive	2194.3326	32.381	67.767	0.000	2130.867	2257.798
cylinders	3494.2375	14.834	235.555	0.000	3465.163	3523.312

```

=====
Omnibus:                97120.996    Durbin-Watson:           1.680
Prob(Omnibus):           0.000    Jarque-Bera (JB):        6333982.669
Skew:                    2.763    Prob(JB):                0.00
Kurtosis:                35.792    Cond. No.:               5.93e+05
=====

```

#### Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 5.93e+05. This might indicate that there are strong multicollinearity or other numerical problems.

## Part 5 Conclusions:

- r2 for 5 models:
  - model 1: 0.312
  - model 2: 0.230
  - model 3: 0.374
  - model 4: 0.553

- model 5: 0.568
- Model 5 has a highest R squared value (0.568), meaning it can explain the dependence variable better compared to other models.

## Concluding Remarks:

- Descriptive data analysis shows popular features or market trends include: white color, four wheel drive, followed by front wheel drive. Ford is the most popular brand.
- Correlation analysis suggested a weak positive correlation between price(dependent variable) and drive type(independent variable), a negative correlation between price(dependent variable) and age and odometer(independent variables).
- Inferential analysis suggested Model 5 (odometer + age + cylinders + drive) is the best model, with a R squared of 0.568, meaning the combination of odometer + age + cylinders + drive can explain 56.8% variance in the listings.

In [ ]:

In [ ]: